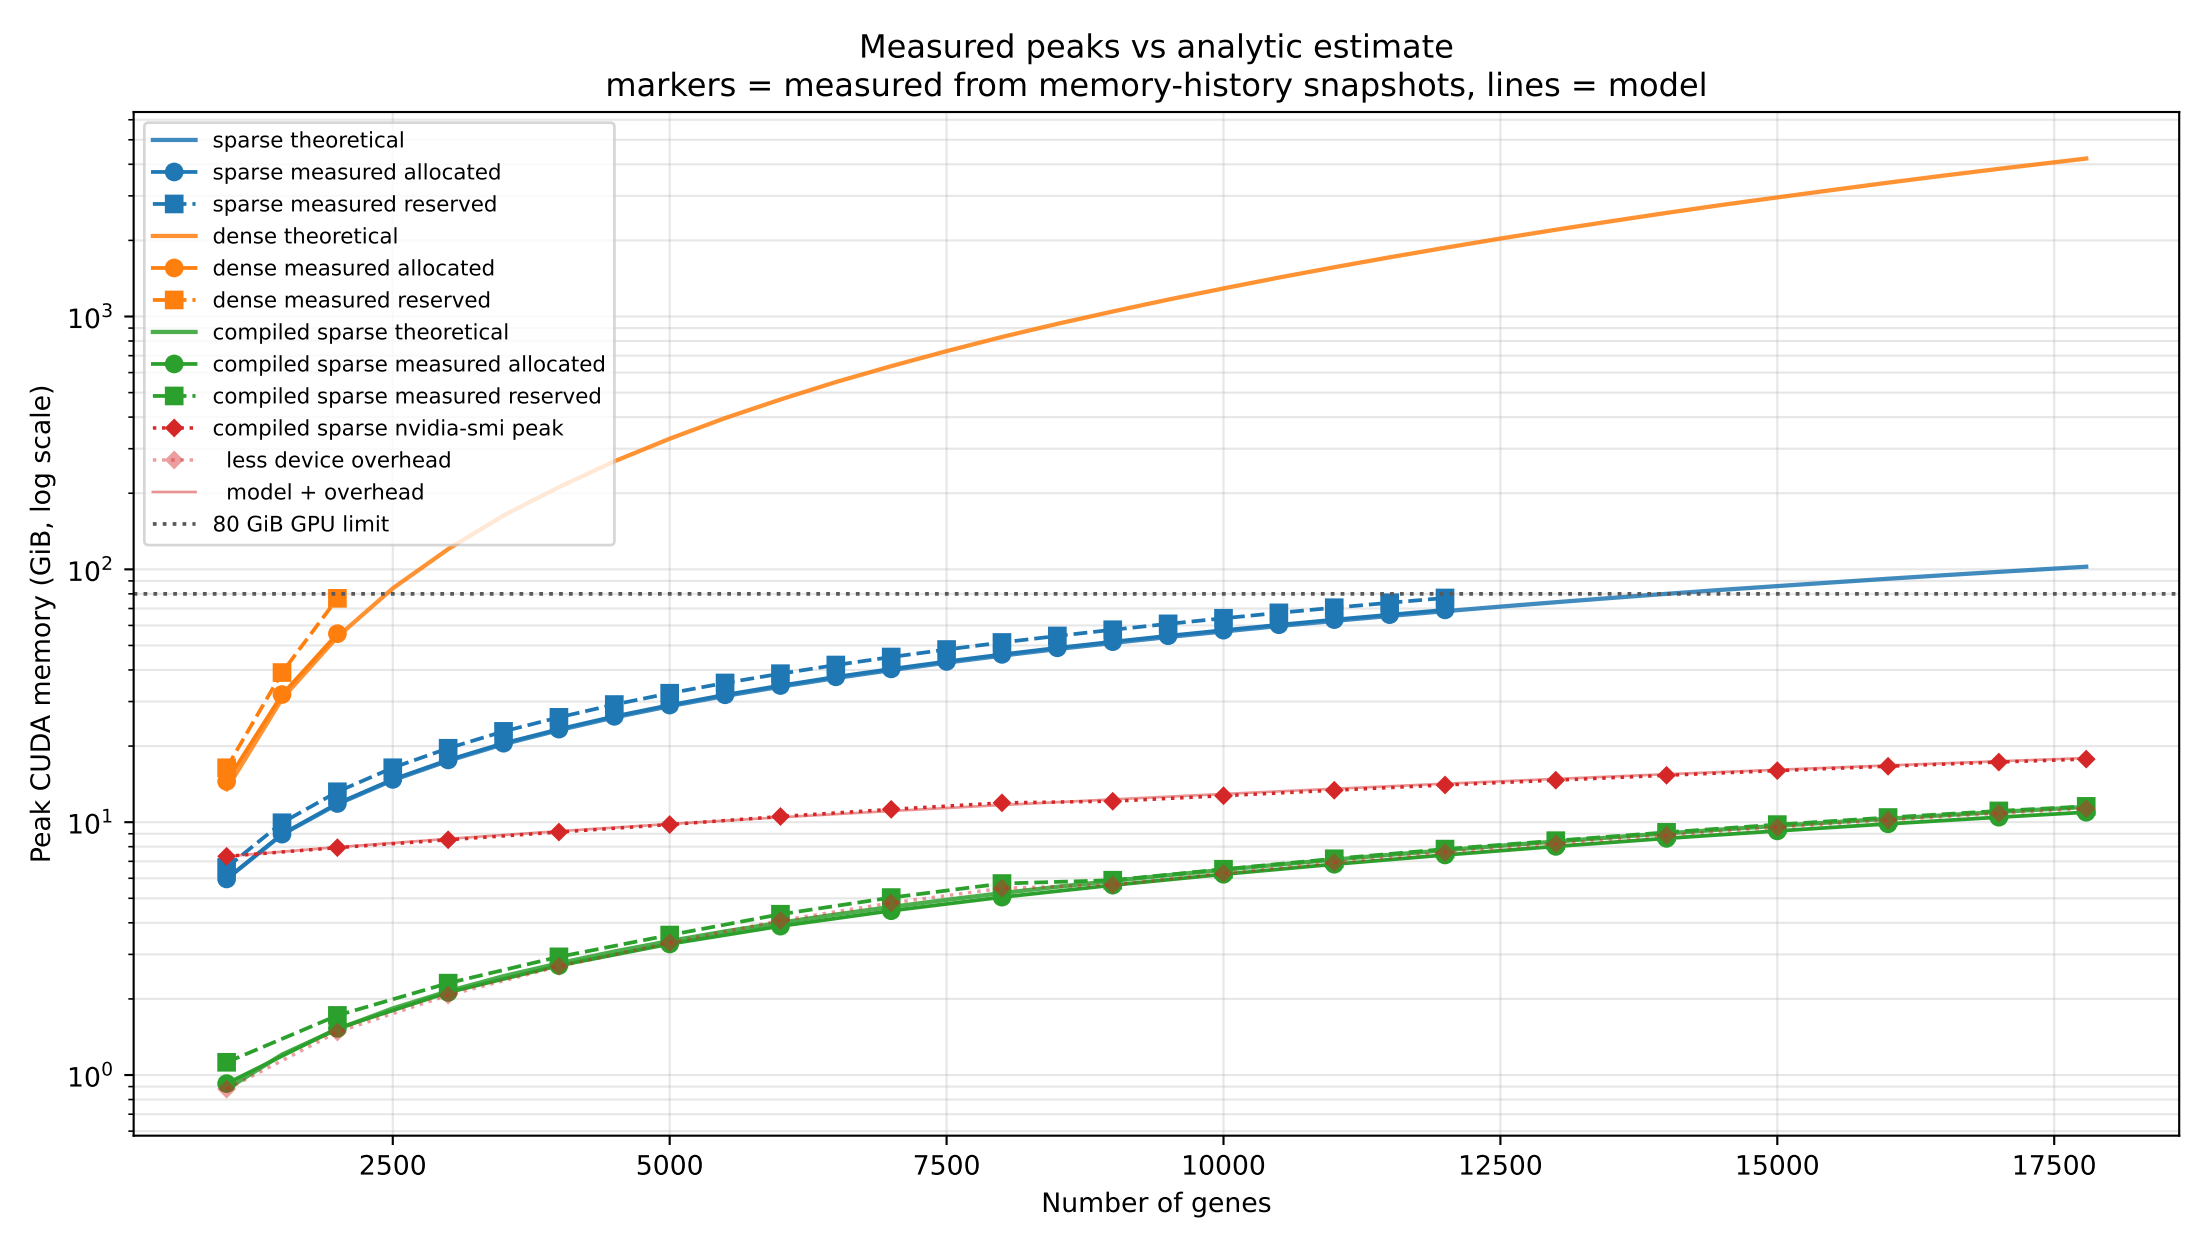
Measured peaks vs analytic estimate
markers = measured from memory-history snapshots, lines = model
sparse theoretical
sparse measured allocated
sparse measured reserved
dense theoretical
dense measured allocated
dense measured reserved
compiled sparse theoretical
compiled sparse measured allocated
compiled sparse measured reserved
compiled sparse nvidia-smi peak
less device overhead
model + overhead
80 GiB GPU limit
Peak CUDA memory (GiB, log scale)
10^3
10^2
10^1
10^0
2500
5000
7500
10000
12500
15000
17500
Number of genes