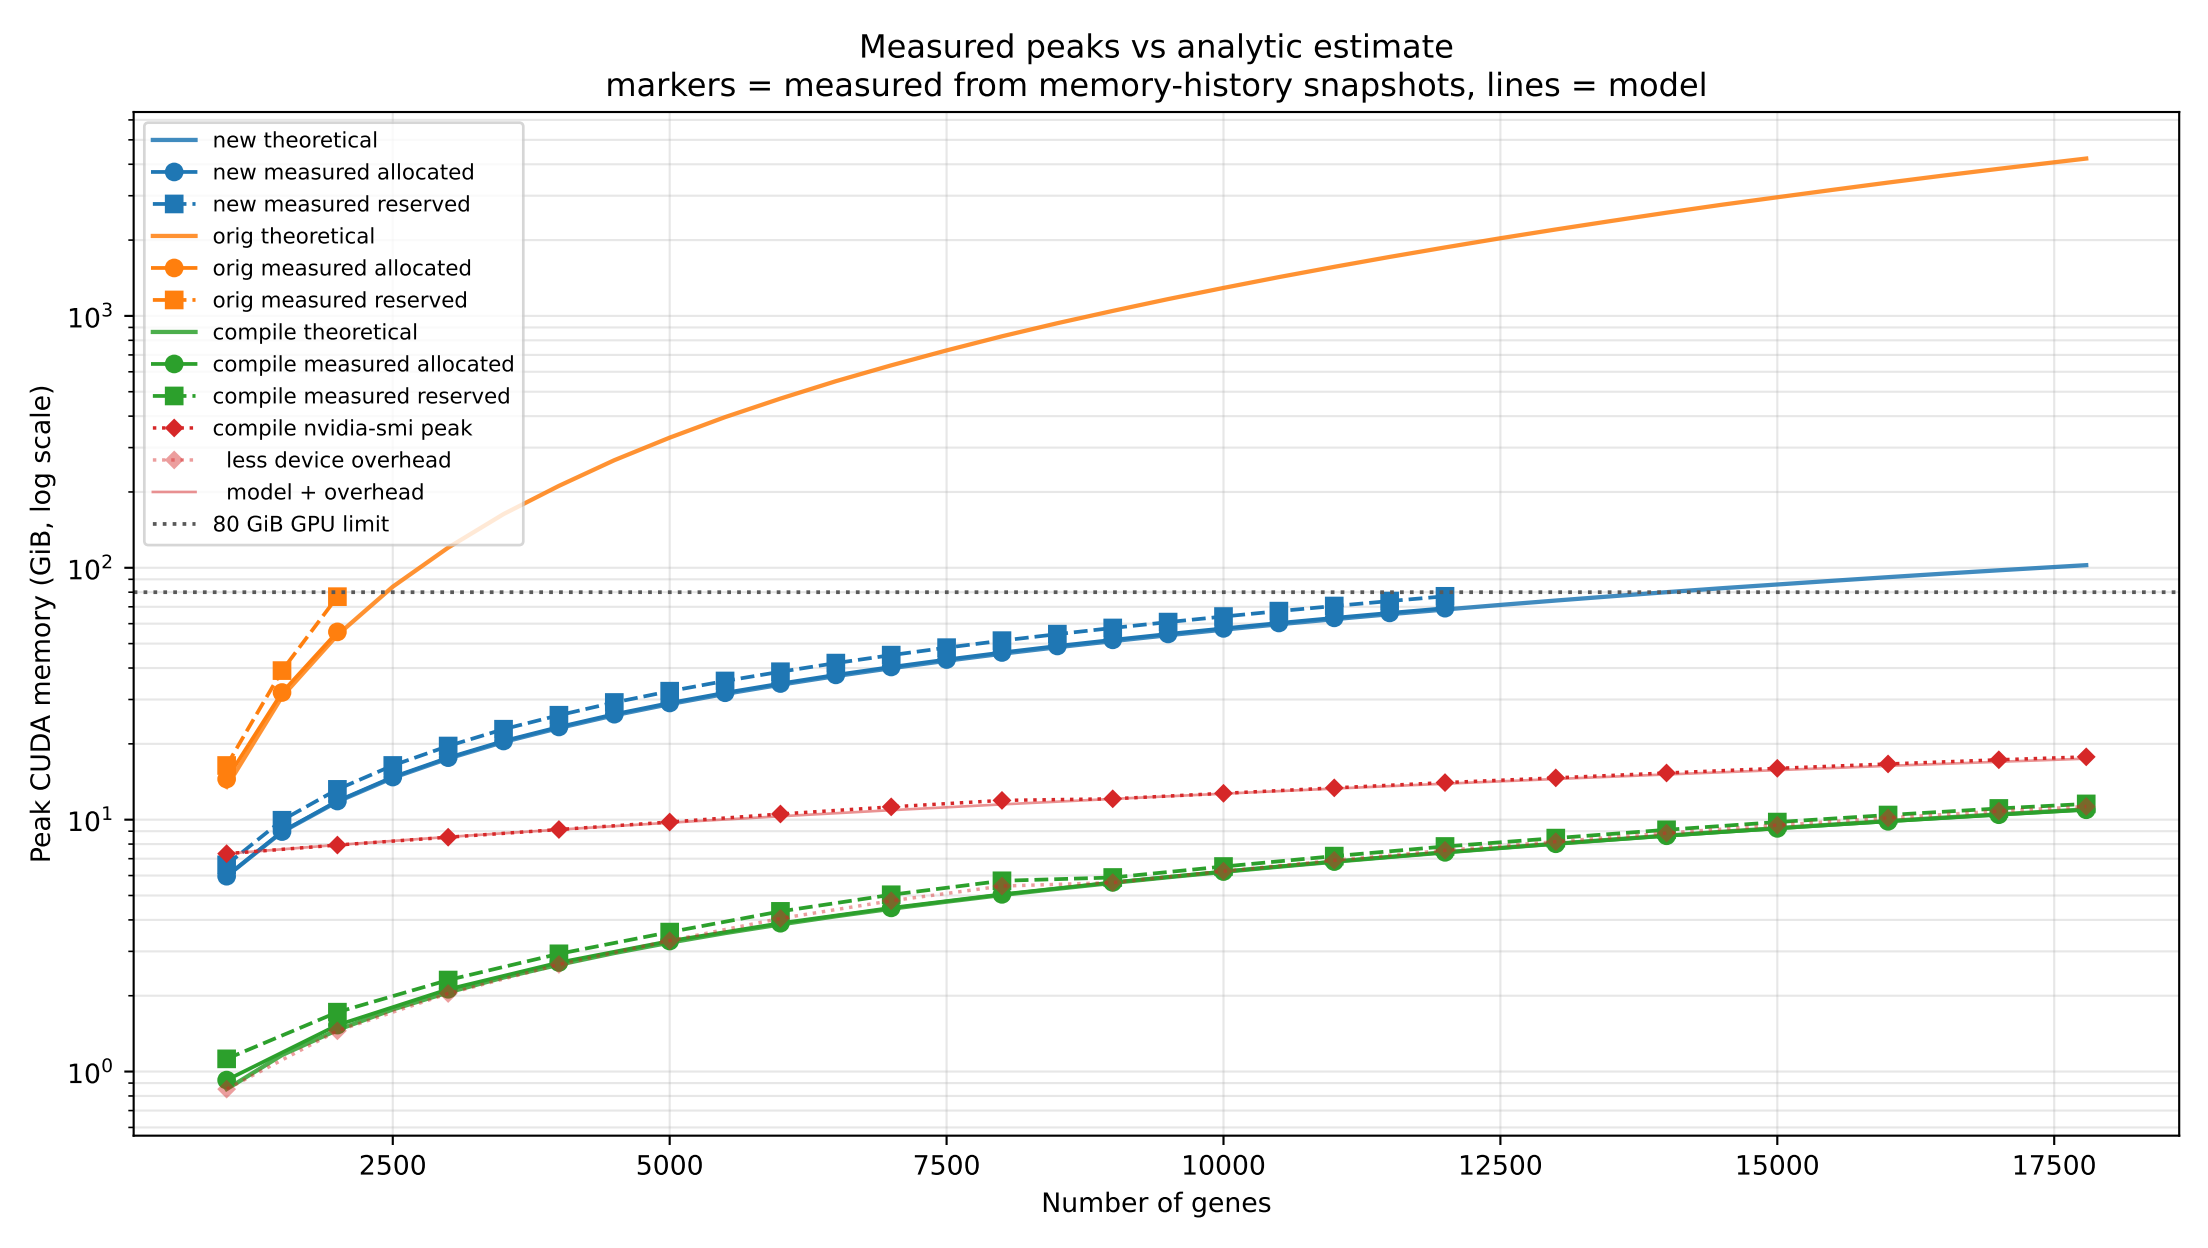
Measured peaks vs analytic estimate
markers = measured from memory-history snapshots, lines = model
new theoretical
new measured allocated
new measured reserved
orig theoretical
orig measured allocated
orig measured reserved
compile theoretical
compile measured allocated
compile measured reserved
compile nvidia-smi peak
less device overhead
model + overhead
80 GiB GPU limit
Peak CUDA memory (GiB, log scale)
10^3
10^2
10^1
10^0
2500
5000
7500
10000
12500
15000
17500
Number of genes