Causal generator peak memory vs TFs per target gene (k)

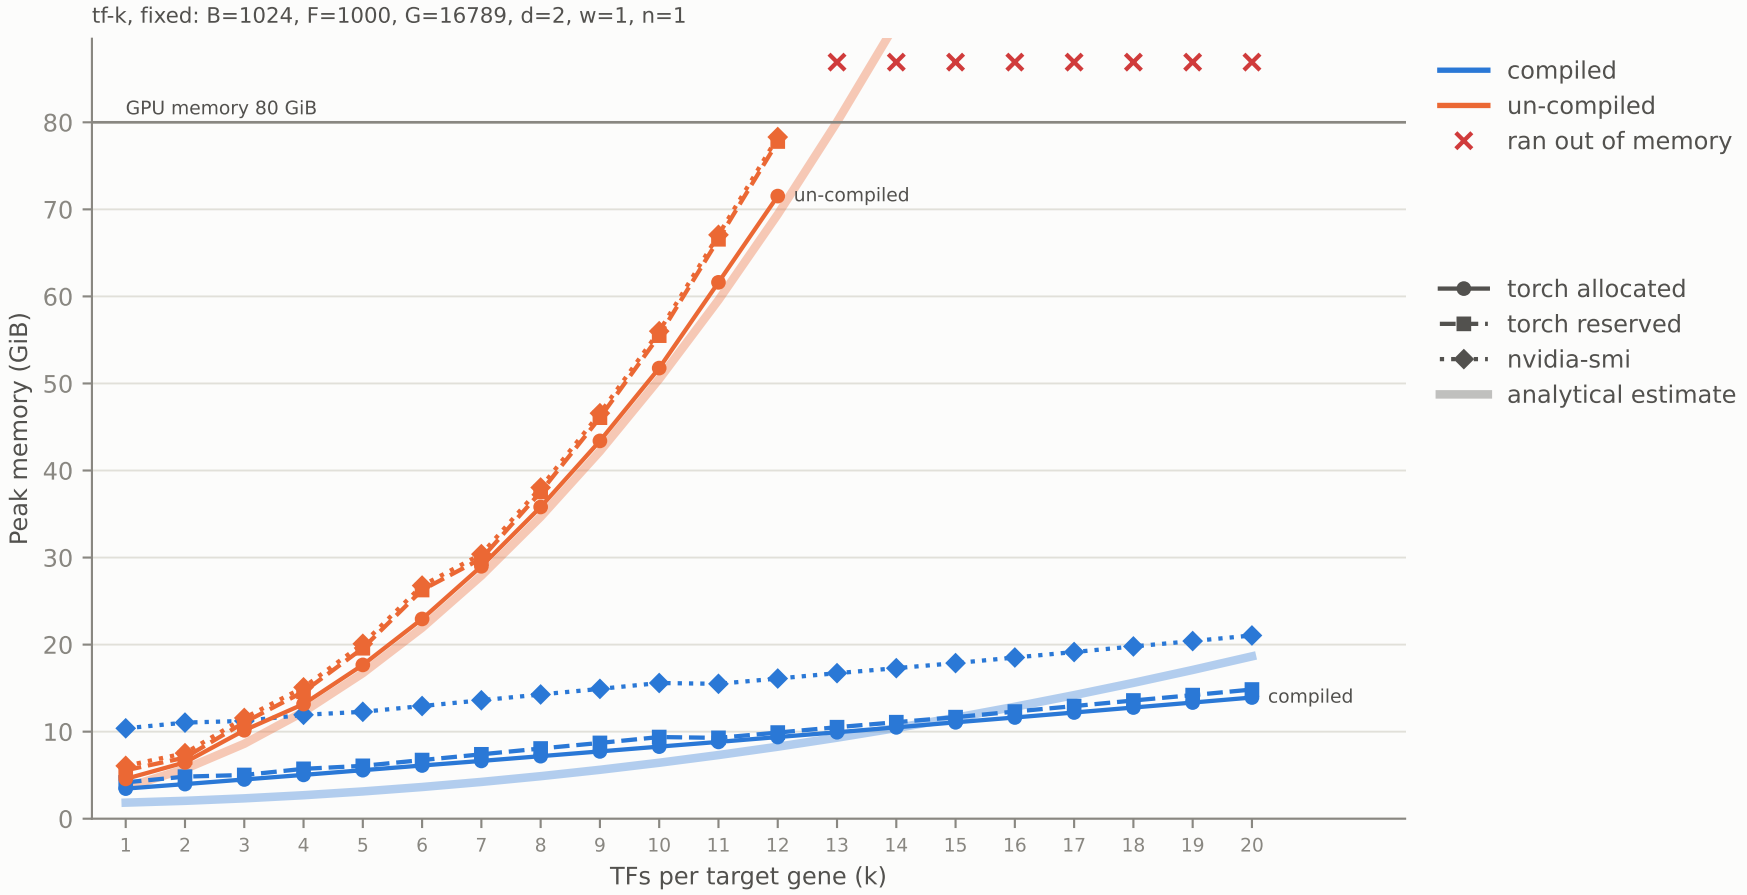

Torch peaks over the whole phase, including the warm-up. The analytical estimate is for the generator step and is drawn beyond the GPU memory, where runs failed.
Compiled nvidia-smi values include the CUDA local-memory pool grown by max-autotune (see compile-overhead/).